

Exploration of Rewards Shaping and Imitation Learning in Chess

04.02.2026

Jakob Lambert-Hartmann

Contents

Introduction	2
Methods	3
Neural Network Architecture	3
Behavior Cloning	4
Tactics and Strategy	4
Puzzle Data	5
Master Games	6
Model-Free Reinforcement Learning	7
REINFORCE Policy Gradient Algorithm	7
Application in our Setting	8
Reward Shaping	8
Results	9
Behavior Cloning	10
Reward Shaping	11
Conclusion	11
Acknowledgement	13
Bibliography	13
Appendix	14

Introduction

Reinforcement learning (RL) is a field of machine learning concerned with sequential decision-making. An agent interacts with an environment by observing a **state** $s_t \in \mathcal{S}$ at time step t , selecting an **action** $a_t \in \mathcal{A}$, and receiving a scalar **reward** $r_t \in \mathbb{R}$. The goal of the agent is to maximize the expected cumulative reward, also called the **return**, $G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k}$, where $\gamma \in [0, 1]$ is a discount factor.

The actions chosen by the agent (e.g. a chess move) are determined by a **policy** $\pi(a | s)$, which defines a probability distribution over actions given a state. A **trajectory** $\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots)$ is generated by interacting with the environment while following the policy π . Winning a game yields a positive reward, losing yields a negative reward, and otherwise the reward is zero. Over the course of training, the agent adapts its policy π in order to maximize the expected return (see Figure 1).




Figure 1: Reinforcement learning: Agent chooses action based on state. Environment reacts to the state and returns a new state and a reward to the agent.

If rewards are sparse, meaning most rewards are zero, it is difficult for an agent to learn a good policy, as there will be too little feedback for the model to learn. Winning a game of chess for instance is a sparse reward. It is difficult for an agent to learn chess with win as only reward, as it is very unlikely for an agent to randomly win a game. To help agents learn policies faster reward shaping introduces further rewards that align with the original goal and help to nudge the model in the right direction [1]. Reward shaping however is challenging and time-consuming. Common difficulties are:

Reward Hacking Agents exploiting unintended loopholes in the reward functions [2]

Misaligned Rewards Rewards are misaligned with the true objective [2]

Difficulty in Evaluation It is difficult to evaluate the effectiveness of a reward function [2]

Imitation learning, on the other hand, is a paradigm in reinforcement learning, in which an agent learns a policy through a set of expert demonstrations, rather than an explicit reward function. Behavior cloning is a specific imitation learning technique, which utilizes super-

vised learning to learn this policy. Using supervised learning circumvents reward shaping, however, expert demonstrations used for behavior cloning are often not uniformly sampled from the state space, which may lead to poor performance [3].

One Idea for mitigating the drawbacks of both reward based reinforcement learning and behavior cloning would be to combine the two approaches. Behavior cloning as a pretraining step could reduce the need for reward shaping, while fine-tuning with reward based reinforcement learning could iron out the sampling bias from the expert demonstrations. In this document we investigate the interaction between pretraining with behavior cloning and fine-tuning with reward shaping using different sets of expert demonstrations and reward functions, which we will apply to the game of chess.

Methods

To investigate the interaction between behavior cloning and reward shaping designed reward functions of varying densities, and selected expert demonstrations with varying sampling bias. We then trained models for all behavior cloning dataset and reward function combinations. For better generalizability of our findings, we performed the same experiment with three neural network architectures.

In the following sections we will describe our choices in model architecture, behavior cloning, and reward shaping in more detail.

Neural Network Architecture

We repeated our experiment using three neural network architectures: a fully connected neural network, a convolutional neural network [4], and a residual network [5].

While the architectures differ in their hidden layers, they share the same input and output representation. The **input** is a one hot encoding of the current position. There are 8×8 squares on a chess board and 12 distinct pieces (black and white version of: pawn, rook, knight, bishop, queen, and king). One simplification we made, is that it's always white to play. To generate a move for black, the position first needs to be mirrored. The reason is that otherwise, we would either need to train two models, or that the model would need to learn the same patterns, once for white and once for black. As chess is a symmetric game, this simplification does not reduce the model's capability.

The network **output** represents a distribution over moves. It is implemented as two tensors (from which move, to which move), with which we compute joint probability distribution. In chess the number of legal moves depends on the position (see Figure 2), as such the move distribution is computed from the set of legal moves. This representation limits the number of parameters for faster training, however it is not without drawbacks:

Limiting Agent Expressiveness In chess, when a Pawn moves up the board it can promote to any piece (e.g. transform into a queen). In our output representation it is only possible to promote to a queen. However, promotions are rare, and almost all pawns are promoted to a queen. This simplification should not affect model performance much.

Inductive Bias The representation of the move distribution as a joint probability introduces inductive bias, that may misalign with chess. This approach implies that the choice of $\text{Square}_{\text{from}}$ and $\text{Square}_{\text{to}}$ are independent. This does not hold true for all positions (e.g. moving the queen to a square may lead to a checkmate, while other pieces would prolong the game).




Figure 2: Red arrows show all the legal moves in both positions. As the number of legal moves changes depending on position, the model masks out illegal moves for to generate actions.

Behavior Cloning

For behavior cloning we chose two datasets (puzzles and master games) with different expert demonstration bias. Models were trained on either dataset for one or 10 epochs. To measure the effect of pretraining, we also had a model that was not subject to pretraining. The following sections will describe both datasets in further detail. Before describing the dataset bias however, we need to introduce some domain specific concepts.

Tactics and Strategy

For an agent to be successful in Chess it needs to master both strategies and tactics. **Tactics** defines moves that achieve short term gains. This may be capturing opponent pieces or checkmating the opponent king. **Strategy** on the other hand defines moves with a long term benefit. This could include positioning a piece to put pressure on the opponents king, which may lead to a checkmate in 20–30 moves. Figure 3 highlights those concepts with two example positions. One key difference between tactical moves and strategic moves, is the

number of viable alternatives. Whereas tactics commonly has only one, or maybe two viable moves in a position, strategy often has a range of viable moves with little difference. This property affects how suitable both aspects are to supervised learning.




Figure 3: The position on the left has a tactical motif, where the black knight can be captured by white, as it is not defended. On the other hand the right position has a strategic motif, where the knight is attacked and has to retreat. There are multiple available squares, but d5 is the best option. There is no immediate gain, however in the long run the knight will be able to exert more pressure on the opponent from this square.

Puzzle Data

The first dataset we used is the Lichess puzzle dataset. [6]. Chess puzzles are positions which contain a tactical motif and a sequence of approximately 1-5 moves to win an advantage. Lichess is a popular online chess website, which generates those puzzles from games played on the website. As such, they contain positions from all levels of play.

While the dataset spans all levels of play, it is still a heavily biased sample of all possible states. Puzzle positions, by nature, contain positions with a single best move. For most chess positions, especially strategic positions, there is no “solution”. Such positions are missing from the dataset. This bias may lead to a network that lacks strategy and is overconfident in tactical themes (see Figure 4)




Figure 4: both positions are similar except for the black bishop on d5. In the position with the Bishop (right) black can win the game in one move by checkmating with the queen on e2. playing queen e2 in the other position will lose the game.

Master Games

The second dataset we used contains games played by Grandmasters [7]. Compared to puzzles, this dataset contains a better balance between tactics and strategy, as it contains entire games. As those are games played by Grandmasters however, they are biased toward higher level play. For Grandmasters it is typical to resign in lost positions, rather than wait for an unavoidable checkmate. The underrepresentation of checkmates may lead to a model that plays well, but is unable to convert a winning position into a win.

Another difficulty with master games is choosing a label for positions in the dataset. The puzzle dataset was well suited to supervised learning, as every position has exactly one winning move. This is not the case for the Grandmaster games (see Figure 5) A few possibilities to deal with this would be to Select one move: the first move encountered is the only label chosen label Predict distribution: aggregate the different positions into distribution of moves and use it as label.

Out of time constraints however, we implemented neither and included copies of the same data-point with different labels. This is reflected in the test accuracy discussed in the next chapter.




Figure 5: The position on the left shows the last position of a game by the two former world champions Garry Kasparov and Veselin Topalov. Black resigns in this completely lost position without waiting for a checkmate. The position on the right shows different first moves played by masters, showing that there is more than one viable move per position.

Model-Free Reinforcement Learning

Model-free reinforcement learning is commonly divided into two main approaches: Q-learning and policy optimization. Q-learning focuses on learning an action-value function $Q_\theta(s, a)$ that approximates the optimal value function Q^* by optimizing the Bellman equation, and it is typically off-policy, meaning it can learn from data generated by any policy; the resulting policy is obtained implicitly by selecting the action that maximizes $Q_\theta(s, a)$. In contrast, policy optimization directly learns a policy function $\pi_\theta(a | s)$ and selects actions without relying on a value function, optimizing the parameters θ via gradient ascent on the expected return; these methods are usually on-policy and use data generated by the current policy only. In terms of trade-offs, policy optimization methods are generally more principled and stable and handle stochastic policies and continuous action spaces more naturally, while Q-learning methods tend to be more sample efficient but can suffer from instability due to bootstrapping and maximization bias. For simplicity, we therefore combine imitation learning with REINFORCE, which is a policy optimization method, in our approach.

REINFORCE Policy Gradient Algorithm

REINFORCE is a foundational policy gradient algorithm in reinforcement learning. It directly optimizes the parameters of a stochastic policy by performing gradient ascent on the expected cumulative return.

Given a parameterized policy $\pi_\theta(a | s)$ that outputs action probabilities, the expected return can be expressed as

$$J(\pi_\theta) = E_{\tau \sim \pi_\theta}[R(\tau)],$$

where τ is a trajectory and $R(\tau)$ is its total discounted reward.

Using the policy gradient theorem [8] this expectation can be rewritten as

$$\nabla_\theta J(\pi_\theta) = E_{\tau \sim \pi_\theta} \left[\sum_{t=0}^T \nabla \log \pi_\theta(a_t | s_t) G_t \right].$$

Here, $G_t = \sum_{t'=t}^T \gamma^{t'-t} r_{t'}$ is the cumulative return-to-go from timestep t .

In practice, REINFORCE estimates this gradient using samples from the policy. This is done by collecting trajectories under the current policy, computing the returns for each action, and taking a step in the direction of $\nabla \log \pi_\theta(a_t | s_t) G_t$.

Application in our Setting

We employ the REINFORCE algorithm to train a neural network-based chess agent through self-play. The policy π_θ is parameterized by a neural network with architecture

$$\theta \in \{\text{fully-connected, CNN, ResNet}\},$$

which outputs a probability distribution over all legal chess moves given a board state.

During training, trajectories are generated by letting the agent play complete chess games. At each timestep t , the policy selects a move a_t given state s_t , and the log-probability $\log \pi_\theta(a_t | s_t)$ of the selected action is recorded. After the game terminates, cumulative discounted returns G_t are computed for every timestep using a combination of reward functions (e.g., checkmate, material advantage).

Reward Shaping

Models trained with reinforcement learning through self play. The models played 16 games before updating based with REINFORCE [9]. This was performed for 1000 iterations, resulting in 16000 games played both as white and as black. After each chosen action a set of reward function was evaluated.

Reward shaping is a technique that provides intermediate rewards to a reinforcement learning agent during training, accelerating convergence without altering the optimal policy. In domains like chess where meaningful outcomes are rare and distant from individual actions, shaping rewards help bridge the credit assignment gap by providing frequent feedback on strategically meaningful decisions.

Training a chess agent purely on win/loss signals is slow — wins happen rarely and far from

the moves that caused them, making it hard for the model to learn. We explore reward shaping by testing three configurations of increasing complexity.

The **sparse** set provides only terminal rewards: wins are rewarded and losses penalized.

The **medium** set adds chess principles that matter at all levels: material advantage, center control, king safety, and winning.

The **dense** set was modeled after chess principles with align with winning on lower levels of play. The rewards included: outer center control, center control, castling, pawn promotion, blunder prevention, material advantage, king safety, checks, and winning.

Across all configurations, bonus rewards are deliberately scaled smaller than the win reward, so the agent learns to pursue chess principles without losing sight of the ultimate goal. This approach lets us measure how much intermediate guidance actually helps the model learn, compared to sparse terminal-only feedback.

Results

To evaluate the trained models, we let them play against each other and recorded the wins, losses, and draws. Every model played 1000 games against all other models with the same architecture. Using tournament chess scoring we aggregated a score (1 point for a win, 0 for a loss, and 0.5 for a draw). The tables in this chapter use a score normalized to a range $[0, 1]$, which can be interpreted similarly to a win percentage. In the following sections we evaluate behavior cloning and imitation learning in isolation in

Table 1 shows the proportion of points gained for all training configurations. There is no dominating pretraining strategy, regarding reinforcement learning however, the best models used the sparse rewards or reinforcement learning altogether.

arch rl	cnn				fc				resnet			
	none	r_0	r_1	r_2	none	r_0	r_1	r_2	none	r_0	r_1	r_2
train												
none	nan	0.266	0.235	0.251	nan	0.289	0.307	0.335	nan	0.408	0.407	0.399
pz_1	0.400	0.478	0.564	0.433	0.507	0.697	0.631	0.623	0.830	0.410	0.410	0.409
pz_10	0.446	0.672	0.515	0.520	0.530	0.637	0.521	0.558	0.925	0.403	0.403	0.403
gm_1	0.396	0.605	0.650	0.572	0.390	0.483	0.523	0.494	0.797	0.411	0.409	0.424
gm_10	0.434	0.714	0.694	0.654	0.418	0.494	0.538	0.526	0.870	0.387	0.398	0.399

Table 1: Expected proportion of points for all models (best models highlighted in aqua). The proportions are means of all matches (e.g. Model using residual network pretrained on puzzle data with no reinforcement learning, wins 92.5% of points against all other residual network models)

Behavior Cloning

For models pretrained with supervised learning it is possible to evaluate them on their test sets. The models trained on puzzle data generalized much better on test (see Table 2). But as mentioned in the discussion on the datasets, this is likely to be caused by the labeling inconsistencies where different moves were played in the same positions.

dataset	epochs	fc	cnn	resnet
puzzle	1	62.0%	69.9%	74.6%
	10	65.8%	73.3%	81.7%
grand masters	1	31.0%	38.5%	41.5%
	10	31.8%	39.6%	43.7%

Table 2: Imitation learning training accuracy

Because of those labeling inconsistencies, game results are a better metric to compare the performance of different pretraining strategy (see Table 3). In contrast to the test accuracy, there is no clear pretraining strategy that dominates other models in play. However, there is a significant improvement in pretraining of any length with any dataset compared to pure reinforcement learning.

dataset	epochs	fc	cnn	resnet
none		0.310	0.251	0.405
puzzle	1	0.615	0.469	0.515
	10	0.562	0.538	0.533
grand masters	1	0.472	0.556	0.510
	10	0.494	0.624	0.513

Table 3: Mean expected proportion of points by pretraining strategy (best model of architecture highlighted).

When only comparing models that did not use reinforcement learning however, pretraining on puzzle data for 10 epochs yields the best performance across all architectures (see Table 4).

dataset	epochs	fc	cnn	resnet
puzzle	1	0.507	0.400	0.830
	10	0.530	0.446	0.925
grand masters	1	0.390	0.396	0.797
	10	0.418	0.434	0.870

Table 4: Expected proportion of points for models not using reinforcement learning. (best model of architecture highlighted).

In conclusion, training on puzzle games for 10 epochs yielded the best results without any reinforcement learning. With reinforcement learning however

Reward Shaping

When comparing the reward functions in models that were not pretrained, there is no clear reward function set, that dominates the others (see Table 5). The fully connected network increases performance as the rewards get denser. The CNN and ResNet however show no trend.

rewards	fc	cnn	resnet	rewards	fc	cnn	resnet
none	0.461	0.419	0.855	r_0	0.289	0.266	0.408
r_0	0.520	0.547	0.404	r_1	0.307	0.235	0.408
r_1	0.504	0.532	0.405	r_2	0.335	0.251	0.399
r_2	0.507	0.486	0.407				

Table 5: Expected points by reward set. Left: across all models. Right: across models without any supervised pretraining.

When considering the reward set across all pretraining strategy however, models with sparse rewards tend to perform slightly better. For the residual network however, there is a significant drop in performance after any type of reinforcement learning (see Table 5). The results highlight the complexity of reward shaping, as our rewards which should have guided the model towards better policies, did not show any increase in performance.

Conclusion

In conclusion, out of the training strategies used, there was no strategy that performed better across all network architectures. However, We were able to show, that reward shaping is indeed difficult, as no rewards or sparse rewards performed better than the complex reward functions.

As our strongest models performed poorly against even a novice chess player, our experiment might yield different results with stronger models. A few possibilities to train stronger models are:

Mixing puzzles and GM games Training on a mix of puzzle and Grandmaster positions could reduce the demonstration bias.

Data preparation on GM games Multiple labels for the same data points prevented the models to generalize well. A strategy to unify the labels could improve the model's performance (e.g. selecting one move or having move distributions as label).

Longer reinforcement learning Only training our models for 1000 batches of reinforcement learning may be too little for any significant improvement.

Smaller reinforcement learning batches Aggregating the loss for 16 games may hinder the model in learning, as it reduces noise too much.

Acknowledgement

Computations for this work were done using resources of the Leipzig University Computing Center.

Bibliography

- [1] Tim Miller, “Reward Shaping.” [Online]. Available: <https://gibberblot.github.io/rl-notes/single-agent/reward-shaping.html>
 - [2] S. Ibrahim, M. Mostafa, A. Jnadi, H. Salloum, and P. Osinenko, “Comprehensive overview of reward engineering and shaping in advancing reinforcement learning applications,” *IEEE Access*, 2024.
 - [3] “Imitation Learning, Stanford University.” [Online]. Available: https://web.stanford.edu/class/cs237b/pdfs/lecture/lecture_10111213.pdf
 - [4] K. O'shea and R. Nash, “An introduction to convolutional neural networks,” *arXiv preprint arXiv:1511.08458*, 2015.
 - [5] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2016, pp. 770–778.
 - [6] “Lichess chess puzzle dataset.” [Online]. Available: <https://www.kaggle.com/datasets/tianmin/lichess-chess-puzzle-dataset>
 - [7] “All GM Chess Games on Chess.com.” [Online]. Available: <https://www.kaggle.com/datasets/dimitrioskourtikakis/gm-games-chesscom>
 - [8] OpenAI, “Spinning Up: Part 3: Intro to Policy Optimization.” [Online]. Available: https://spinningup.openai.com/en/latest/spinningup/rl_intro3.html
 - [9] J. Zhang, J. Kim, B. O'Donoghue, and S. P. Boyd, “Sample Efficient Reinforcement Learning with REINFORCE,” *CoRR*, 2020, [Online]. Available: <https://arxiv.org/abs/2010.11364>
-

Appendix
